

# lab1

January 21, 2025

## 1 Digit recognition with a CNN

Code to initialize Tensorflow 2.0 in Colab

```
[1]: from __future__ import absolute_import, division, print_function, \
      unicode_literals
      # %tensorflow_version 2.x
      import tensorflow as tf
      %load_ext tensorboard
      import datetime
      import numpy as np
      import matplotlib.pyplot as plt
```

```
2025-01-21 22:30:47.615947: I tensorflow/core/util/port.cc:153] oneDNN custom
operations are on. You may see slightly different numerical results due to
floating-point round-off errors from different computation orders. To turn them
off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-01-21 22:30:47.616890: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:32]
Could not find cuda drivers on your machine, GPU will not be used.
2025-01-21 22:30:47.623159: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:32]
Could not find cuda drivers on your machine, GPU will not be used.
2025-01-21 22:30:47.642046: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
E0000 00:00:1737495047.673600 178167 cuda_dnn.cc:8310] Unable to register cuDNN
factory: Attempting to register factory for plugin cuDNN when one has already
been registered
E0000 00:00:1737495047.683101 178167 cuda_blas.cc:1418] Unable to register
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has
already been registered
2025-01-21 22:30:47.714122: I tensorflow/core/platform/cpu_feature_guard.cc:210]
This TensorFlow binary is optimized to use available CPU instructions in
performance-critical operations.
To enable the following instructions: AVX2 AVX512F AVX512_VNNI FMA, in other
operations, rebuild TensorFlow with the appropriate compiler flags.
```

Import the MNIST dataset. The default loader will return tensors for the train/test partitions of the images and the labels.

```
[2]: mnist = tf.keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train[:,:,:,:np.newaxis]/255.0 # our images are 4 dimensional
↳ (NumImages, Height, Width, Channels) - these images are grayscale so we only
↳ have 1 channel
x_test = x_test[:,:,:,:np.newaxis]/255.0
```

[TODO] Check the size of the loaded tensors

```
[3]: print(f'Dimension of X_train: {x_train.shape}') # 6000 images used as training
↳ set, each with size 28x28 and only one channel cause they are in greyscale
print(f'Dimension of X_test: {x_test.shape}') # 1000 images used as test set
print(f'Dimension of Y_train: {y_train.shape}') # 6000 labels
print(f'Dimension of Y_test: {y_test.shape}') # 1000 labels
```

Dimension of X\_train: (60000, 28, 28, 1)

Dimension of X\_test: (10000, 28, 28, 1)

Dimension of Y\_train: (60000,)

Dimension of Y\_test: (10000,)

Prepare Keras callback for Tensorboard

```
[4]: logdir = "logs/scalars/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
%tensorboard --logdir logs
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=logdir,
↳ update_freq='batch')
```

<IPython.core.display.HTML object>

[TODO] Define a Keras Sequential model with the convolutional neural network

```
[5]: # we are performing images classification, so given a dataset of labeled images
↳ (supervised learning) we want that our model understand how to extract
↳ features and then use these to predict the class of the image
# for the process of features extraction we use convolutional layers, that
↳ performs F different filters on the same image in order to obtain a map
↳ associated to each filter. These maps correspond to some features
# extracted from the images. These process of features extraction is called
↳ convolutional encoding, cause through the convolutional layers we are
↳ extracting more and more details from the image and we are reducing
# the size of the image. In our case we use 'same' padding, this means that the
↳ on the input image are added new pixels in the border, so the convolution
↳ can return an image with the same size. This new pixels can
# be added equals to 0 (zero padding), or with the same values of the real
↳ borders (symmetric padding), or with the same values just over the border
↳ (reflection padding). We could decide to not perform any padding,
```

# so we give to the convolutional layer the original image, this method is  
 ↳ called valid padding and the returned output will be the image reduced of 2  
 ↳ pixels on each dimension.

# While we are going deeper in the NN, the convolutional layers should extract  
 ↳ more features cause we want to extract more and more complex features. Note  
 ↳ that a pixel in the output image of convolutional layer is

# is evaluated on a perceptive field of the size of the kernel used by the  
 ↳ convolution, so while we are going deeper each pixel contains the  
 ↳ information of more pixels.

# Between the layers we are also performing Batch Normalization, this is  
 ↳ important cause it normalize (by standardization) the input values by  
 ↳ considering the statistics (mean and standard deviation) of a batch

# of the data. This is important cause it reduce the effect of too low value  
 ↳ that could bring to vanishing gradient, and too high value that could bring  
 ↳ to exploding gradient.

# The we use an activation function, that recives the input values and perform  
 ↳ on them the ReLu activation function. The activation function perform an  
 ↳ operation of the input. Is important that this function is

# non linear, cause if it was linear it will perform the same operations on  
 ↳ each input, and so we loose the advantages of having multiple layers cause  
 ↳ they are all performing the same action.

# We can decide between sigmoid, ReLu and Leaky ReLu for applying the non  
 ↳ linear activation function. Sigmoid function is the worst, cause all the low  
 ↳ input are forced to 0 and all the high input are forced to 1,

# this cause saturation of the outputs and so vanishing gradient. ReLu is a  
 ↳ linear function ( $y=x$ ) on positive inputs, but it force all the negative  
 ↳ inputs to 0. Leaky ReLu is the same but it doesn't force

# the negative input to 0 but it set them with a linear function ( $y=x/10$ ).

# Between the convolutional layers we are also performing Pooling, these  
 ↳ operations reduces the dimension of the image (not of the features) by  
 ↳ selecting a specific value (max, min, avarege) on the kernel.

# For applying the classification we use softmax function that returns a pdf  
 ↳ between all the 10 possible classes. Before doing so, softmax function need  
 ↳ a vector as input, so we have to flat the matrix into a

# vector by using Flatten layer. From the pdf of the softmax we read the class  
 ↳ with the higher probability. This is how a convolutional encoder works for  
 ↳ performing image classification.

# Before deciding the class i've added a new layer that perform dropout. This  
 ↳ will remove randomly 50% of the neurons from the trained network and then  
 ↳ try to train the other neurons to perform the same.

# This method reduces the dimension of the neural network but it increases the  
 ↳ values of the parameters. Parameters increases of  $1/(1-p)$ , where  $p$  is the  
 ↳ percentage of neurons freed. Having high values of the

# parameters could be not good, cause high weights made the model very sen  
 ↳ sistive on input channges. So we could also perform regularization to give  
 ↳ penalty on the cost function for high weights.

```

# I have tested that the model with the regularization perform worst than the
↳ model without it, so i decided to remove regularization.
# Note that the NN is divided into two steps, before it applies features
↳ extraction and then it performs decision making on those.
model = tf.keras.models.Sequential([
    # tf.keras.layers.Conv2D(32, (3,3), padding='same', kernel_regularizer=tf.
↳ keras.regularizers.L2(0.01)), # extracts basics features from the image by
↳ using a convolutional kernel (3x3) and 32 filters, so we get 32 feature maps.
↳ With the regularization the cost function receive a penalty that follows
↳ the square norm of the weights.
    tf.keras.layers.Conv2D(32, (3,3), padding='same'),
    tf.keras.layers.BatchNormalization(), # this layer normalize the output by
↳ applying standardization by using the statistics evaluated on a batch
↳ extracted from the training dataset. Used to avoid exploding and vanishing
↳ gradient
    tf.keras.layers.ReLU(), # non linear activation function. All the negative
↳ input are forced to 0, the positive ones are setted following a linear
↳ function y=x
    tf.keras.layers.MaxPooling2D((2,2)), # halve the size of the image by
↳ taking the maximum value from a 2x2 square
    # tf.keras.layers.Conv2D(64, (3,3), padding='same', kernel_regularizer=tf.
↳ keras.regularizers.l2(0.01)), # we increase the number of filters to 64,
↳ so we get 64 feature maps. This layer will extract more complex features
    tf.keras.layers.Conv2D(32, (3,3), padding='same'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.ReLU(),
    tf.keras.layers.MaxPooling2D((2,2)), # halve the size of the image by
↳ taking the maximum value from a 2x2 square
    # tf.keras.layers.Conv2D(128, (3,3), padding='same', kernel_regularizer=tf.
↳ keras.regularizers.l2(0.01)), # we extract again more complex features
    tf.keras.layers.Conv2D(32, (3,3), padding='same'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.ReLU(),
    tf.keras.layers.Flatten(), # we flatten the 3D tensor to a 1D tensor,
↳ because Dense layer only accepts 1D tensor
    tf.keras.layers.Dropout(0.5), # we use dropout to avoid overfitting, this
↳ layer randomly set 0 some neurons, then the NN is forced to train with the
↳ remeining ones how to work as previous
    tf.keras.layers.Dense(10, activation='softmax') # softmax activation
↳ function is used to get the probabilities of each class, so the output is a
↳ vector of 10 probabilities (one for each class) summed to 1
])

```

W0000 00:00:1737495058.276980 178167 gpu\_device.cc:2344] Cannot dlopen some GPU libraries. Please make sure the missing libraries mentioned above are installed properly if you would like to use GPU. Follow the guide at <https://www.tensorflow.org/install/gpu> for how to download and setup the

required libraries for your platform.  
Skipping registering GPU devices...

[TODO] Compile the Keras model: specify the optimization algorithm, the loss function and the test metric

```
[6]: # Train a NN means trying to find the best weights and biases (parameters of
    ↪ the NN) that can explain in the best way possible the relations between
    ↪ inputs and outputs. For doing so we must evaluate we amount
    # of error made by our model, so we can understand how far we are from the
    ↪ desired result. This evaluation is done through the cost function, that
    ↪ could be the MSE (mean square error) or the cross entropy
    # evaluated between the desired output and the one that our NN returns. After
    ↪ having evaluated the amount error through the cost function, we have to
    ↪ understand how to change the parameters in order to reduce the
    # error. This could be done by evaluating the gradient, so we derive the cost
    ↪ function on each parameter and we obtain a vector in the space of the
    ↪ function that indicates to us the direction where we have to move
    # for reducing the error. Our goal is to arrive to the minimum point of this
    ↪ function by exploiting the gradient, so we evaluate the cost function and
    ↪ then we update the parameters with the gradient until we reach
    # a convergence. This algorithm is called gradient descending, and it ensures
    ↪ that we could reach a local minimum point, but not that it is the global
    ↪ minimum. We could start from random parameters and see where
    # the algorithm brings us. Then we can try again with other random parameter
    ↪ and see if the gradient descend converge on better minimum. This is the way
    ↪ how we can increase the probability of reaching the global
    # minimum. Note that the gradient gives us only the direction where we have to
    ↪ move, but not the amount of step that we should perform. For finding this
    ↪ value we should evaluate the second derivative, but this mean
    # to evaluate (number of parameters)2 derivative of the cost function, too
    ↪ complex in system with many parameters as the NN. So what we can do is to
    ↪ estimate the amount of step through a learning rate.
    # A too big step can cause jumping over minimum point and so guide on the wrong
    ↪ direction the algorithm, instead a too short step could slow down the
    ↪ convergence of the algorithm.
    # We know that the cost function after having applied the step should be lower
    ↪ that the actual value:  $C(w + \Delta w) - C(w) = -\eta \cdot \nabla_w (C)$ , where  $\eta$  is the learning
    ↪ rate.
    # At each iteration gradient descending must read the entire dataset for
    ↪ evaluating the cost function and then update the parameters with the
    ↪ gradient. An optimization is the Stochastic Gradient Descending, that
    # performs the same actions but it doesn't load the entire dataset but it
    ↪ extract from it a batch and then it doesn't put it back. So at each
    ↪ iteration the cost function is evaluated on the extracted batch, and
```

```

# the size of the training set decreases of the batch size. The number of
↳ epochs represent the number of times that the algorithm should see the
↳ entire training set, so the total number of iterations performed
# by this algorithm is  $(n / b) * e$ , where  $n$  is the size of the training set,  $b$ 
↳ is the size of the batch and  $e$  is the number of epochs.
# Adam oprimizer is an optimization of the stochastic gradient descent, that
↳ changes the learnig rate through the iterations by considering the
↳ statistics of the gradient.
lr = 0.01 # learning rate
model.compile(optimizer = tf.keras.optimizers.Adam(lr), loss =
↳ 'sparse_categorical_crossentropy', metrics=['accuracy'])

```

[TODO] Train the Keras model

```

[7]: model.fit(x_train, y_train, batch_size=128, epochs=5,
↳ callbacks=[tensorboard_callback])

```

```

Epoch 1/5
469/469          41s 78ms/step -
accuracy: 0.8818 - loss: 0.4391
Epoch 2/5
469/469          34s 72ms/step -
accuracy: 0.9778 - loss: 0.0691
Epoch 3/5
469/469          32s 68ms/step -
accuracy: 0.9824 - loss: 0.0551
Epoch 4/5
469/469          31s 66ms/step -
accuracy: 0.9841 - loss: 0.0504
Epoch 5/5
469/469          32s 67ms/step -
accuracy: 0.9865 - loss: 0.0437

```

```

[7]: <keras.src.callbacks.history.History at 0x79aa808ff6a0>

```

[TODO] Print model summary

```

[8]: model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 32)	320
batch_normalization (BatchNormalization)	(None, 28, 28, 32)	128

re_lu (ReLU)	(None, 28, 28, 32)	0
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 14, 14, 32)	9,248
batch_normalization_1 (BatchNormalization)	(None, 14, 14, 32)	128
re_lu_1 (ReLU)	(None, 14, 14, 32)	0
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 32)	0
conv2d_2 (Conv2D)	(None, 7, 7, 32)	9,248
batch_normalization_2 (BatchNormalization)	(None, 7, 7, 32)	128
re_lu_2 (ReLU)	(None, 7, 7, 32)	0
flatten (Flatten)	(None, 1568)	0
dropout (Dropout)	(None, 1568)	0
dense (Dense)	(None, 10)	15,690

Total params: 104,288 (407.38 KB)

Trainable params: 34,698 (135.54 KB)

Non-trainable params: 192 (768.00 B)

Optimizer params: 69,398 (271.09 KB)

[TODO] Test the Keras model by computing the accuracy the whole test set

```
[9]: model.evaluate(x_test, y_test)
```

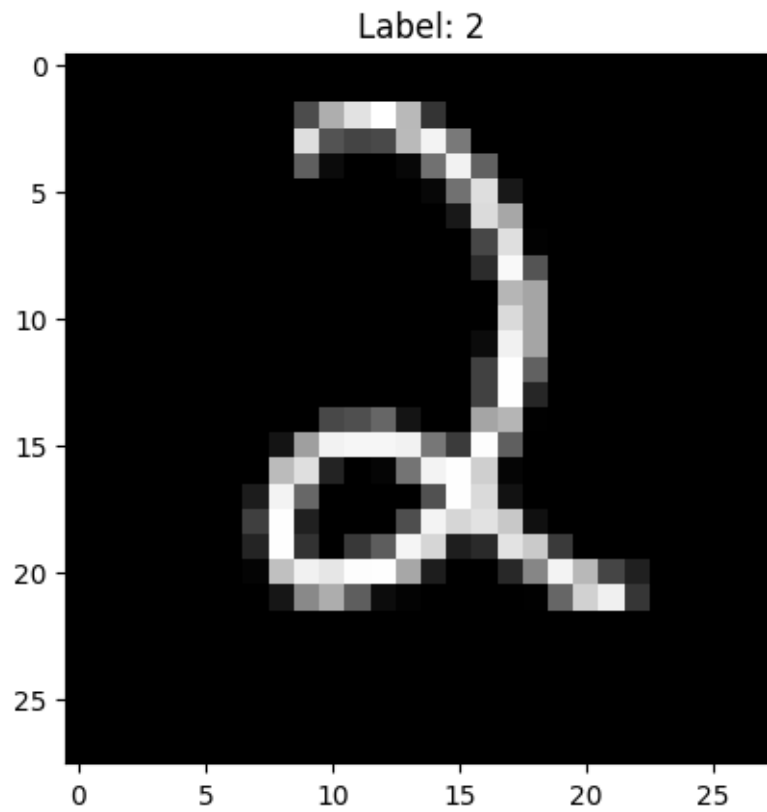
```
313/313          3s 7ms/step -
accuracy: 0.9712 - loss: 0.1088
```

```
[9]: [0.09016285091638565, 0.9763000011444092]
```

[TODO] Visualize test image number 47 and the prediction from the neural network

```
[10]: plt.imshow(x_test[47].reshape(28, 28), cmap='gray')
plt.title(f'Label: {y_test[47]}')
plt.show()

y_pred = model.predict(x_test[47][np.newaxis, :, :, :])
print(f'Predicted label: {np.argmax(y_pred)}')
```



1/1            0s 221ms/step  
Predicted label: 2